

Some title*

Extended Abstract[†]

Maximilian Bachl
Institute of Telecommunications
Vienna, Austria
maximilian.bachl@tuwien.ac.at

Tanja Zseby
Institute of Telecommunications
Vienna, Austria
tanja.zseby@tuwien.ac.at

Joachim Fabini
Institute of Telecommunications
Vienna, Austria
joachim.fabini@tuwien.ac.at

ABSTRACT

This paper provides a sample of a \LaTeX document which conforms, somewhat loosely, to the formatting guidelines for ACM SIG Proceedings.¹

CCS CONCEPTS

• **Computer systems organization** → **Embedded systems**; *Redundancy*; Robotics; • **Networks** → Network reliability;

KEYWORDS

Put some comma-separated keywords here

ACM Reference Format:

Maximilian Bachl, Tanja Zseby, and Joachim Fabini. 2017. Some title: Extended Abstract. In *Proceedings of ACM Woodstock conference (WOODSTOCK'97)*. ACM, New York, NY, USA, 5 pages. https://doi.org/10.475/123_4

1 INTRODUCTION

In recent years, there have been several approaches to make congestion control learnable, most notably *Remy* [3], which learns an optimum congestion control w.r.t. a specific utility function for specified network scenarios in an off-line fashion. Furthermore there is *Performance-oriented Congestion Control* (PCC) [1], which repeatedly tries different rates during periods of approximately one round-trip time (RTT) and chooses the better rate w.r.t. a specific utility function.

We make use of a new Reinforcement Learning (RL) based approach that can learn a congestion control algorithm on-line and is able to perform actions on a per acknowledgement basis.

1.1 Actor Critic Learning

To this end we propose *Partial Action Learning* (PAL) (see 2.1) which is a modification of the Actor Critic framework for neural networks proposed by Mnih et al. [2], which we outline in this section.

There are two Artificial Neural Networks (ANNs), the Actor Network and the Value Network (the Critic part in the abbreviation stands for the Value Network). Given a state, the Actor Network

outputs what it deems to be the optimum action to perform in that certain state. The Value Network estimates what long-term reward can be expected in this state. So an action is considered good if it achieved a long-term reward that is higher than the long-term reward expected by the Value Network and it is considered bad if the reward was lower than expected. The long-term reward is implemented as a moving average of future rewards. So if a high reward can be achieved right now this is more favorable than if it can be achieved in the future. However it can also be beneficial to get a low reward now and instead get a very large one in the future.

1.1.1 Value Network. The Value Network outputs the expected long-term reward $V(s_t; \theta_v)$ given a state s_t at time step t and the parameters (neural network weights) of the value network θ_v .

With r_t being the reward that was received at time t and γ being the roll-off factor, which stands for the influence that future reward has on the moving average (commonly set to 0.99), we define the expected long-term reward as

$$R_t = \left(\left(\sum_{i=0}^{k-1} \gamma^i r_{t+i} \right) + \gamma^k V(s_{t+k}; \theta_v) \right) (1 - \gamma),$$

where k is upper-bounded by t_{\max} (t_{\max} is a fixed hyperparameter that indicates how many rewards should be received before updating the neural network). So R_t is simply a moving average of rewards at time step t . However, usual moving averages take into account values from the past while this one uses values from the future. In practice this means that we collect t_{\max} rewards, compute the expected long-term reward starting at each time step, update the neural network and start the same procedure again.

The loss function, which the value network tries to minimize at each time step, is the square of the difference of the actual long-term reward received and the expected long-term reward

$$l_{v,t} = (R_t - V(s_t; \theta_v))^2.$$

1.1.2 Actor Network. The Actor Network outputs a probability distribution from which the action a_t at time step t is randomly sampled. We use the mean μ and the standard deviation σ , to parametrize a normal distribution. The main idea is that the network learns to output the right mean at the right time step to maximize the future reward and that it uses the standard deviation to try out new actions, which could yield a better than expected reward.

With v_t being the value that the value network estimated as the future reward given the current state s_t at time step t (v_t is just an abbreviation for $V(s_t; \theta_v)$) and θ_a the parameters of the actor network and with β being a factor that specifies the importance of the entropy H , π designating the probability density function, meaning that $\pi(a_t | s_t; \theta_a)$ is the value of the probability density function of taking action a_t in state s_t with the current weights

*Produces the permission block, and copyright information

[†]The full version of the author's guide is available as `acmart.pdf` document

¹This is an abstract footnote

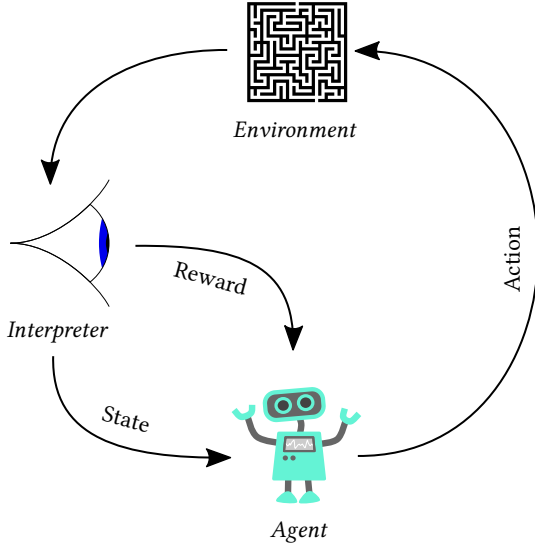


Figure 1: The classical Reinforcement Learning approach.^a

^aadapted from https://upload.wikimedia.org/wikipedia/commons/1/1b/Reinforcement_learning_diagram.svg

of the actor network θ_a , we define the loss that the actor network aims to minimize as follows:

$$l_{a,t} = -\log(\pi(a_t | s_t; \theta_a)) (R_t - v_t) - \beta H(\pi(s_t; \theta_a)).$$

2 METHOD

2.1 Partial Action Learning

The key difference between PAL and previous approaches to Reinforcement Learning is that in classical Reinforcement Learning, an action is always followed by a reward and a reward is always followed by an action. In our proposed concept, however, it is possible to take new actions while previous actions haven't received their rewards yet.

Another major difference in PAL is that one action generates a number of partial actions (≥ 0) (see Figure 2). Each partial action generates feedback upon interacting with the environment. Upon receiving feedback for a partial action, the agent determines the current state and triggers a new action. When all feedbacks of one action were received, the agent combines them to form the reward and updates the value and actor networks.

In Algorithm 1 we show the code that runs in each of the agents (in the congestion control scenario, one agent corresponds to a sender). It is possible to have several agents which share a set of weights but one can also use separate weights for each agent, which is more realistic in case of congestion control in the Internet, as different senders cannot easily share a set of neural network weights over the Internet.

2.2 Congestion Control specifics

The motivation for PAL is that classical Reinforcement Learning assumes that a reward follows an action and vice-versa (see Figure 1).

Algorithm 1 Partial Action Learning – pseudocode for each agent. It is possible that the agents share the global weights $\theta_{a,g}$ and $\theta_{v,g}$, which can be reasonable when learning off-line. Otherwise an individual copy of them is kept by each agent. All weights are initialized randomly in the beginning. LSTM cells are used in both the actor and the value network.

```

1: loop
2:    $l_{actions} \leftarrow []$ 
3:    $l_{states} \leftarrow []$ 
4:    $l_{values} \leftarrow []$ 
5:    $l_{rewards} \leftarrow []$ 
6:    $l_{estimatedValues} \leftarrow []$ 
7:    $l_{snapshots} \leftarrow []$ 
8:    $l_{hiddenStates} \leftarrow []$ 
9:    $t \leftarrow 0$ 
10:   $s_0 \leftarrow initialState()$ 
11:   $h_a \leftarrow initialHiddenState()$ 
12:   $h_v \leftarrow initialHiddenState()$ 
13:   $\theta_a \leftarrow \theta_{a,g}$ 
14:   $\theta_v \leftarrow \theta_{v,g}$ 
15:  repeat
16:     $l_{estimatedValues} \cdot append(V(s_t; \theta_v))$ 
17:    if  $t \bmod t_{max} = 0$  then
18:       $\theta_a \leftarrow \theta_{a,g}$ 
19:       $\theta_v \leftarrow \theta_{v,g}$ 
20:       $l_{snapshots} \cdot append((\theta_a, \theta_v))$ 
21:       $l_{hiddenStates} \cdot append((h_a, h_v))$ 
22:    end if
23:     $l_{states} \cdot append(s_t)$ 
24:    Sample  $a_t$  from the
      Actor Network's probability distribution
25:     $l_{actions} \cdot append(a_t)$ 
26:     $l_{values} \cdot append(V(s_t; \theta_v))$ 
27:     $l_{partialActions, t} \leftarrow partialActions(a_t)$ 
28:    for all partial actions  $a_{p,i,t}$  in  $l_{partialActions, t}$  do
29:      Take partial action  $a_{p,i,t}$ 
30:    end for
31:    Wait for the next feedback  $r_{p,j,t'}$ 
      where  $t' \leq t$  and  $0 \leq j \leq \#(l_{partialActions, t'})$ 
32:    if all feedback of  $a_{p,i,t}$  was received then
33:       $r_{t'} \leftarrow$  reward w.r.t all feedback  $r_{p,k,t'}$ 
      where  $0 \leq k \leq \#(l_{partialActions, t'})$ 
34:       $l_{rewards} \cdot append(r_{t'})$ 
35:      if  $\#(l_{rewards}) > t_{max}$  or the episode is over then
36:        COMPUTEGRADIENTS
37:      end if
38:    end if
39:    Generate  $s_{t+1}$  using  $r_{p,i,t'}$ 
40:     $t \leftarrow t + 1$ 
41:  until reaching the end the episode
42: end loop

```

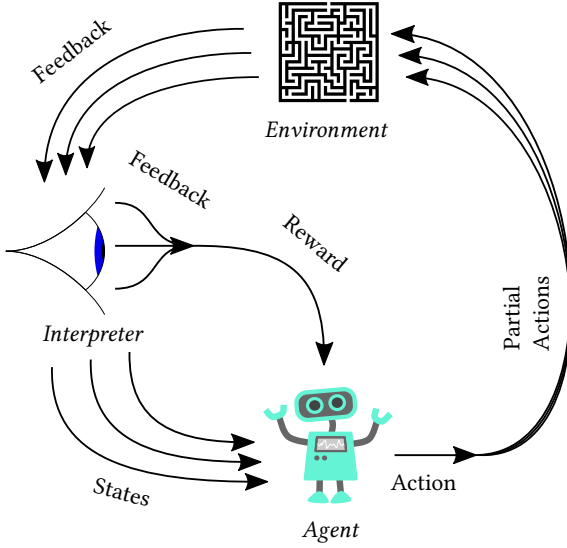


Figure 2: Partial Action Learning: An action consists of zero or more partial actions which trigger feedback upon interacting with the environment. Each feedback updates the state. The value and actor networks are updated upon receiving all feedback of one action.^a

^aadapted from https://upload.wikimedia.org/wikipedia/commons/1/1b/Reinforcement_learning_diagram.svg

Algorithm 2 Partial Action Learning – procedure which computes and applies the gradients.

```

1: function COMPUTEGRADIENTS
2:    $t_{\text{end}} = \min(t_{\text{max}}, \#(l_{\text{rewards}}))$ 
3:    $\theta_{\text{backup}} \leftarrow (\theta_a, \theta_v)$ 
4:    $h_{\text{backup}} \leftarrow (h_a, h_v)$ 
5:    $\theta_a, \theta_v \leftarrow l_{\text{snapshots}}[0]$ 
6:    $h_a, h_v \leftarrow l_{\text{hiddenStates}}[0]$ 
7:    $R_{i+1} \leftarrow \text{last element of } l_{\text{estimatedValues}}$ 
8:   for  $i \leftarrow t_{\text{end}} - 1, 0$  do
9:      $R_i \leftarrow (r_i + \gamma R_{i+1})(1 - \gamma)$ 
10:     $a \leftarrow l_{\text{actions}}[i]$ 
11:     $s \leftarrow l_{\text{states}}[i]$ 
12:     $v \leftarrow l_{\text{values}}[i]$ 
13:     $d\theta_a \leftarrow -\frac{\partial \log(\pi(a | s; \theta_a))(R_i - v)}{\partial \theta_a}$ 
14:     $d\theta_v \leftarrow \frac{\partial (R_i - v)}{\partial \theta_v}$ 
15:     $\theta_{a,g} \leftarrow \theta_{a,g} + d\theta_a$ 
16:     $\theta_{v,g} \leftarrow \theta_{v,g} + d\theta_v$ 
17:   end for
18:   Remove first  $t_{\text{end}}$  elements from
      $l_{\text{actions}}, l_{\text{states}}, l_{\text{values}}, l_{\text{rewards}}, l_{\text{estimatedValues}}$ 
19:   Remove the first element from  $l_{\text{snapshots}}$ 
20:   Remove the first element from  $l_{\text{hiddenStates}}$ 
21:    $\theta_a, \theta_v = \theta_{\text{backup}}$ 
22:    $h_a, h_v = h_{\text{backup}}$ 
23: end function

```

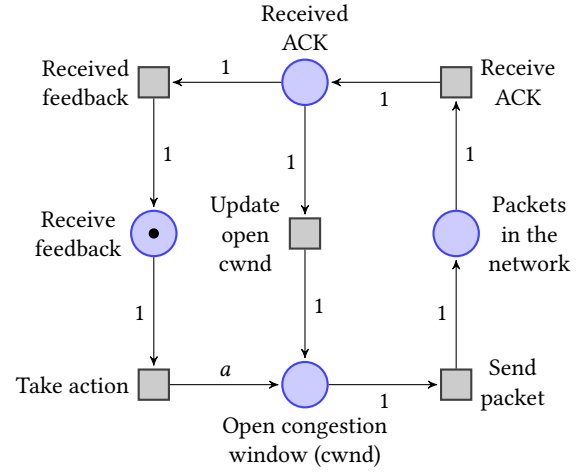


Figure 3: A petri net describing the congestion control mechanism. We start with one token in the state *Received feedback*, which means that we can take an action right in the beginning. The concept is the following: If there is at least one token in the *Open congestion window*, send a packet. The token goes to the network and after some time the acknowledgement (or timeout if the packet gets lost) for the packet is received in the *Received ACK* state. From here one token goes to the *Open congestion window* meaning that when we get an ACK (or timeout) we can send another packet (The *Open congestion window* signifies the congestion window minus the packets that are currently unacknowledged.). Furthermore the state *Received feedback* receives a token upon receiving an ACK/timeout. *Take action* adds a tokens to the *Open congestion window*, where a is a real number (possibly also negative); thus in each state there can also be a real number of tokens (e.g. 2.34 tokens are possible in the *Open congestion window*). However, we define that the congestion window can never be smaller than 1 (This means that the sum of all tokens in the *Open congestion window*, *Received ACK* and *Packets in the network* states can never be smaller than 1).

However, in the case of congestion control, it is desirable to perform a new action without having received a reward for the previous action. For example, imagine that we receive two acknowledgements directly after one other. For each of these two acknowledgements an action has to be performed but it does not seem feasible for the second action to wait until the first action has received a reward as it takes one RTT until an acknowledgement (a reward) is received. Thus the Asynchronous Actor Critic framework as described by [2] cannot be applied to congestion control as it assumes that actions and rewards are synchronized and so we have to use Partial Action Learning (see subsection 2.1). We describe the overall workings of PAL for congestion control using a petri net (see Figure 3).

To use PAL for congestion control we first have to define the correct semantics for this specific use case and we have to explicitly state how the state, reward etc. are defined. Furthermore, we have to define how the actor and value network explicitly work in case

of congestion control. In the following a time step t corresponds to the reception of an acknowledgement. The beginning of the flow, before any packet is sent, corresponds to time step 0.

s_t The state describes the current “congestion state”. The following features are included in it:

- the round-trip time of the last packet
- the current congestion window
- the time between the last two packets that were sent
- the time between the last two packets that were received
- the number of packets that were lost since the last acknowledgement was received

Each time an acknowledgement is received, the state is updated and the actor network is asked for the next action.

a_t Based on a given state and the history of previous states (because we use LSTM cells and thus can also consider previous states), the actor network returns an action a_t , which is a real number that stands for the change of the congestion window.

r_t The reward is a tuple of at least one reward metric. For each reward metric there is also an output of the value network that predicts the expected long-term average of this reward metric given the current state.

We actually use the following types of reward:

$r_{\text{packet},t}$ is the number of the packets that the sender sent during time step t and that were not lost (so they were acknowledged at some point by the receiver).

$r_{\text{delay},t}$ is the sum of the round trip times of the packets that the sender sent and that were not lost.

$r_{\text{duration},t}$ is the sum of the time between receiving the last packet and receiving this packet (“inter-receive time”) for the packets that the sender sent and that were not lost.

$r_{\text{lost},t}$ is the sum of packets that were lost between receiving the last packet and receiving this packet.

The overall structure of the neural network for both the value and the actor network is depicted in Figure 4.

2.2.1 Value Network. In the case of congestion control, the loss function $l_{v,t}$ of the value network is actually the sum of the squares of the difference for each of the expected long-term averages and the empirically found averages for each reward metric:

$$l_{v,t} = \left(R_{\text{packet},t} - V_{\text{packet}}(s_t; \theta_v)\right)^2 + \left(R_{\text{delay},t} - V_{\text{delay}}(s_t; \theta_v)\right)^2 + \left(R_{\text{duration},t} - V_{\text{duration}}(s_t; \theta_v)\right)^2 + \left(R_{\text{lost},t} - V_{\text{lost}}(s_t; \theta_v)\right)^2$$

2.2.2 Actor Network. The actor network outputs two parameters: The mean of a normal distribution μ and its standard deviation σ .

Each time an action a_t is requested, it is sampled from the current normal distribution defined by the parameters μ and σ .

Then the window is incremented by a_t ; however, the window can never be smaller than 1. At the beginning of a flow the window starts with 1 as well.

With H being the entropy, the actor network minimizes the loss

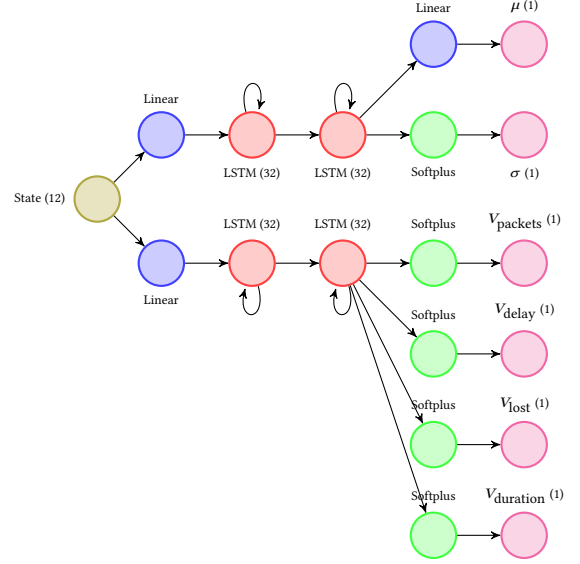


Figure 4: An overview of the complete neural network being used. Ocher stands for the input, blue for linear layers, red for LSTM layers, green for linear layers with a softplus activation function and purple for the outputs of the neural network. The numbers in parentheses next to each label stand for the number of width of that layer (i.e. the number of neurons except for the inputs and outputs).

$$l_{a,t} = -\log(\pi(a_t | s_t; \theta_a)) \left(U_{\text{measured},t} - U_{\text{expected},t} \right) - \beta H(\pi(s_t; \theta_a))$$

where U can be any utility function defined based on some of the reward metrics. In the above formulation, one considers how much better the actual experienced Utility was compared to the expected one. The Utility is formed from the previously defined expected long-term values.

2.3 Utility function

As the utility function one can choose an arbitrary function based on the reward metrics previously defined. A safe choice is to use PCC’s [1] reward function as its convergence has been previously proven. It is defined as follows:

$$U_t = r_{\text{received},t} \text{Sigmoid}_{\alpha} \left(\frac{r_{\text{loss},t}}{r_{\text{sent},t}} - 0.05 \right) + r_{\text{loss},t}$$

where

$$\text{Sigmoid}_{\alpha}(x) = \frac{1}{1 + e^{\alpha x}},$$

r_{received} is the actual throughput (in packets per second), r_{loss} is the loss rate (in packets per second) and r_{sent} is the sending rate (in packets per second). Thus $\frac{r_{\text{loss}}}{r_{\text{sent}}}$ is the ratio of data getting lost compared to those being sent, which is a number between 0 and 1.

Intuitively it means that one takes the actual throughput (packets that get received by the receiver and do not get lost) times the sigmoid function of the loss ratio minus 0.05 plus the loss rate. The sigmoid function essentially acts as a cutoff threshold. As soon as the loss rate rises above 5%, the Sigmoid function decreases very quickly and thus diminishes the Utility.

It is also possible to set the cutoff threshold to 0, which means that losing packets is heavily punished; or one can increase the threshold to allow for more loss.

$r_{\text{received}, t}$ can be calculated as $\frac{R_{\text{packet}, t}}{R_{\text{duration}, t}}$, $r_{\text{loss}, t}$ can be calculated as $\frac{R_{\text{lost}, t}}{R_{\text{duration}, t}}$ and $r_{\text{sent}, t}$ can be calculated as $\frac{R_{\text{packet}, t} + R_{\text{lost}, t}}{R_{\text{duration}, t}}$.

Actually, we assume equally sized packets for all of the above methodology, however, it would also easily be possible to consider the number of bytes per packet as well: One would have to add an additional reward metric of the average number of bytes per packet and as well add one more output to the value network.

2.4 Example

In this section, an example of the overall procedure is outlined.

- (1) Receive an ACK
- (2) Update the internal state to reflect the information that this ACK provides (for example, the state consists of the round trip time experienced by this packet, the time since we receive the last ACK etc.)
- (3) Sample an increase for the congestion window from the Actor Network (e.g. 0.24) and add it to the current congestion window.
- (4)(a) If this was the last ACK, which completes a previous action (e.g. a previous action resulted in 3 packets being sent; so if this is the ACK for the third (and last) packet of this previous action, we got all the ACKs to compute the reward), combine all partial rewards to form the actual reward for that action.
- (b) If we now have t_{max} rewards (set to 20) then we also update the neural network.

REFERENCES

- [1] Mo Dong, Qingxi Li, Doron Zarchy, Philip Brighten Godfrey, and Michael Schapira. 2015. PCC: Re-architecting Congestion Control for Consistent High Performance.. In *NSDI*. 395–408.
- [2] Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi Mirza, Alex Graves, Timothy Lillicrap, Tim Harley, David Silver, and Koray Kavukcuoglu. 2016. Asynchronous methods for deep reinforcement learning. In *International Conference on Machine Learning*. 1928–1937.
- [3] Keith Winstein and Hari Balakrishnan. 2013. Tcp ex machina: Computer-generated congestion control. In *ACM SIGCOMM Computer Communication Review*, Vol. 43. ACM, 123–134.